

## Assignment - 2

---

# GCP VM Auto-Scaling & Security Configuration

---

Submitted By

Abhilash Agarwal  
M23CSA502

<https://github.com/AbhilashAgarwalIITJ/GCP-VM>



# 1 Introduction

## Project Overview

This report documents the implementation of a Google Cloud Platform (GCP) Virtual Machine (VM) with auto-scaling policies and robust security measures. The project aims to create a scalable and secure infrastructure that can adapt to changing workload demands while maintaining strict access controls.

## 1.1 Project Objectives

- Create and configure a VM instance with Apache web server
- Implement auto-scaling using Managed Instance Groups (MIG)
- Configure security measures including firewall rules and IAM roles
- Test the implementation using various tools and techniques

## 1.2 Technical Stack

- Google Cloud Platform (GCP)
- Bash scripts for automation
- Apache HTTP Server
- Testing tools: curl, nmap, stress
- Optional: Terraform for Infrastructure-as-Code

# 2 Architecture Overview

The architecture consists of a Managed Instance Group (MIG) with an underlying instance template that defines VM configurations. Auto-scaling policies are applied to the MIG to dynamically adjust the number of instances based on CPU utilization. Security is enforced through firewall rules that restrict access to specific IP addresses and IAM roles that provide controlled access to GCP resources.

## 2.1 System Components

- VM Instance: Base infrastructure running Apache web server
- Instance Template: Blueprint for creating identical VM instances
- Managed Instance Group: Group of VMs managed as a single entity
- Auto-scaling Policy: Rules for scaling instances based on CPU utilization
- Firewall Rules: Network security rules to control traffic
- IAM Roles: Identity and access management controls

## 3 System Architecture Diagram

### Architecture Overview

The following diagram illustrates the complete architecture of our GCP VM Auto-Scaling and Security Configuration. It shows how different components interact within the Google Cloud Platform environment, including the Managed Instance Group with auto-scaling capabilities, security layers implemented through firewall rules and IAM roles, and the interaction with external clients.

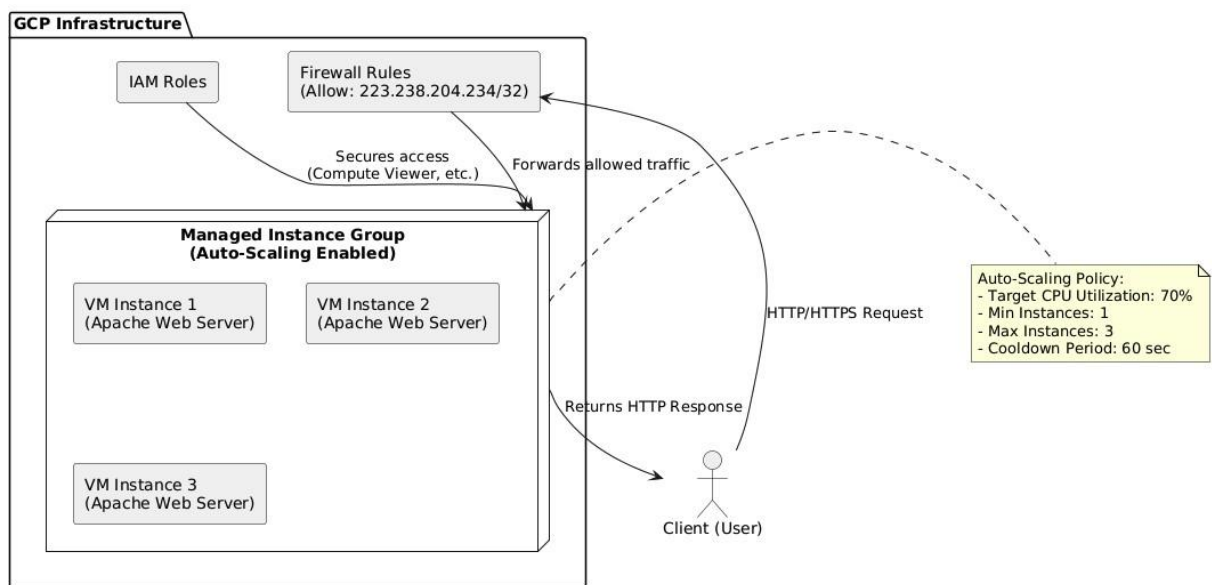


Figure 1: System Architecture Diagram of GCP VM Auto-Scaling & Security Implementation

### 3.1 Key Components

- **Managed Instance Group (MIG):** Central component that manages VM instances and implements auto-scaling based on CPU utilization.
- **Instance Template:** Defines the configuration of VM instances within the MIG, including machine type, operating system, and startup scripts.
- **Auto-Scaling Policy:** Configures when and how the MIG scales, with parameters including target CPU utilization (70%), minimum (1) and maximum (3) instances, and cooldown period (60 seconds).
- **Firewall Rules:** Control network traffic to the VM instances, allowing HTTP/HTTPS access only from authorized IP addresses (223.238.204.234/32).
- **IAM Roles:** Manage user access permissions to GCP resources, with the Compute Viewer role providing read-only access to compute resources.
- **Apache Web Server:** Runs on each VM instance to serve HTTP requests and provide content to users.

### 3.2 Component Interactions

The architecture functions as follows:

1. External users send HTTP/HTTPS requests to the VM instances.
2. GCP Firewall evaluates requests based on the source IP address, allowing only traffic from the authorized IP range (223.238.204.234/32).
3. Authorized requests are forwarded to the Managed Instance Group.
4. The MIG distributes requests among available VM instances running Apache web servers.
5. When CPU utilization exceeds 70%, the auto-scaling policy triggers the creation of additional VM instances (up to a maximum of 3).
6. When CPU utilization decreases, excess VM instances are removed after the 60-second cooldown period (maintaining at least 1 instance).
7. IAM roles ensure that only authorized personnel can view or manage the compute resources.

### 3.3 Security Layers

#### Multi-layered Security Approach

The architecture implements security at multiple levels:

- **Network Level:** Firewall rules restrict HTTP/HTTPS access to specific IP addresses.
- **Identity Level:** IAM roles control which users can access GCP resources and what actions they can perform.
- **Instance Level:** VM instances are configured with minimal required services and regularly updated.
- **Infrastructure Level:** Auto-scaling ensures availability during traffic spikes while preventing resource exhaustion.

## 4 Implementation

### 4.1 Creating a VM Instance

To create a VM instance with Apache installed, we use the following script:

```

1      #!/bin/bash
2      # Create a single VM instance on GCP with Apache installed
3
4      # Set your project ID and zone
5      PROJECT_ID="gcp-vm-assignment"
6      ZONE="us-central1-a"
7      INSTANCE_NAME="gcp-vm01"
8      MACHINE_TYPE="e2-medium"
9      IMAGE_FAMILY="ubuntu-2004-lts"
10     IMAGE_PROJECT="ubuntu-os-cloud"
11
12     # Set the active project
13     gcloud config set project $PROJECT_ID
14
15     # Create the VM instance with HTTP and HTTPS tags
16     gcloud compute instances create $INSTANCE_NAME \
17         --zone=$ZONE \
18         --machine-type=$MACHINE_TYPE \
19         --image-family=$IMAGE_FAMILY \
20         --image-project=$IMAGE_PROJECT \
21         --tags=http-server,https-server \
22         --metadata=startup-script='#!/bin/
23     bash apt-get update
24     apt-get install -y apache2
25     systemctl start apache2
26     systemctl enable apache2'
27
28     echo "VM instance $INSTANCE_NAME created with Apache installed."
```

Listing 1: Create VM with Apache installed

### VM Configuration Details

- Ubuntu 20.04 LTS as the operating system
- e2-medium machine type (2 vCPUs, 4 GB memory)
- HTTP and HTTPS network tags for firewall rule targeting
- A startup script that installs and configures Apache

The screenshot shows a terminal window titled "SSH-in-browser" with a toolbar containing "UPLOAD FILE", "DOWNLOAD FILE", and other icons. The terminal output displays system updates, the status of the Apache service, and its logs.

```
Expanded Security Maintenance for Applications is not enabled.

29 updates can be applied immediately.
25 of these updates are standard security updates.
To see these additional updates run: apt list --upgradable

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

New release '22.04.5 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Sun Mar  2 13:06:20 2025 from 35.235.244.33
m23csa502@assignment-2-gcp-vm:~$ sudo systemctl status apache2
● apache2.service - The Apache HTTP Server
   Loaded: loaded (/lib/systemd/system/apache2.service; enabled; vendor preset: enabled)
   Active: active (running) since Sun 2025-03-02 13:05:00 UTC; 1h 45min ago
     Docs: https://httpd.apache.org/docs/2.4/
   Main PID: 2274 (apache2)
    Tasks: 55 (limit: 4680)
   Memory: 5.1M
   CGroup: /system.slice/apache2.service
           └─2274 /usr/sbin/apache2 -k start
             └─2276 /usr/sbin/apache2 -k start
               └─2277 /usr/sbin/apache2 -k start

Mar 02 13:05:00 assignment-2-gcp-vm systemd[1]: Starting The Apache HTTP Server...
Mar 02 13:05:00 assignment-2-gcp-vm systemd[1]: Started The Apache HTTP Server.
m23csa502@assignment-2-gcp-vm:~$
```

Figure 2: Apache Installation on VM Instance

## 4.2 Configuring Auto-Scaling

Auto-scaling is implemented in two steps:

**Step 1:** Creating an instance template

**Step 2:** Creating a Managed Instance Group with auto-scaling policies

### 4.2.1 Creating an Instance Template

```
1 #!/bin/bash
2 # Create an instance template for auto-scaling
3
4 PROJECT_ID="gcp-vm-assignment"
5 TEMPLATE_NAME="assignment-2-template"
6 MACHINE_TYPE="e2-medium"
7 IMAGE_FAMILY="ubuntu-2004-lts"
8 IMAGE_PROJECT="ubuntu-os-cloud"
9
10 gcloud config set project $PROJECT_ID
11
12 gcloud compute instance-templates create $TEMPLATE_NAME \
   --machine-type=$MACHINE_TYPE \
   --image-family=$IMAGE_FAMILY \
   --image-project=$IMAGE_PROJECT \
```

13  
14  
15  
16  
17

```

18 apt-get update
19 apt-get install -y apache2
20 systemctl start apache2
21 systemctl enable apache2'
22
23 echo "Instance template $TEMPLATE_NAME created successfully."

```

Listing 2: Create an instance template for auto-scaling

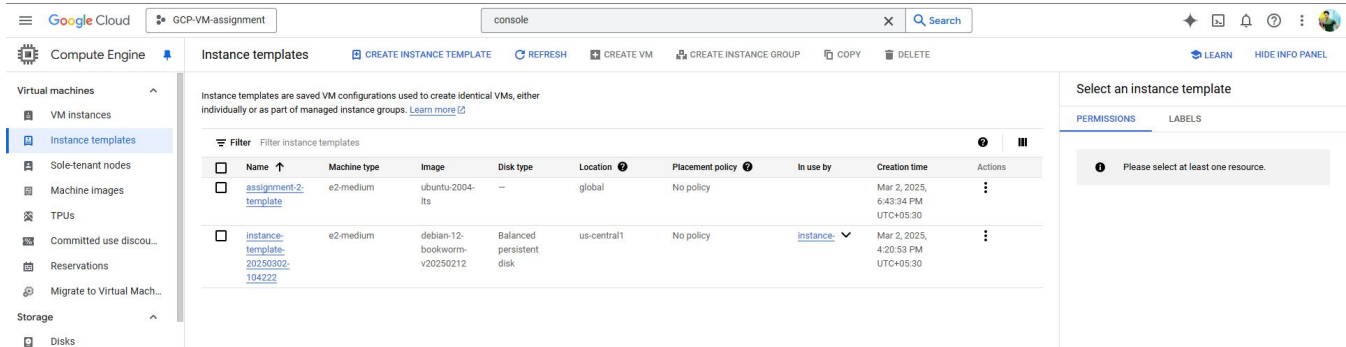


Figure 3: Created Instance Template in GCP Console

#### 4.2.2 Creating a Managed Instance Group with Auto-Scaling

```

1 #!/bin/bash
2 # Create a Managed Instance Group (MIG) with auto-scaling
3
4 PROJECT_ID="gcp-vm-assignment"
5 ZONE="us-central1-a"
6 MIG_NAME="assignment-2-mig"
7 TEMPLATE_NAME="assignment-2-template"
8
9 gcloud config set project $PROJECT_ID
10
11 # Create the Managed Instance Group
12 gcloud compute instance-groups managed create $MIG_NAME \
13   --base-instance-name=assignment-2-mig-instance \
14   --template=$TEMPLATE_NAME \
15   --size=1 \
16   --zone=$ZONE
17
18 # Configure auto-scaling: target CPU utilization 70%, min 1 instance, max 3
19   instances, cooldown 60s
20 gcloud compute instance-groups managed set-autoscaling $MIG_NAME \
21   --zone=$ZONE \
22   --cool-down-period=60 \
23   --max-num-replicas=3 \
24   --min-num-replicas=1 \
25   --target-cpu-utilization=0.7
26
27 echo "Managed Instance Group $MIG_NAME created with auto-scaling configured."
28

```



```

27 echo "Auto-scaling parameters:"
28 echo "  - Target CPU utilization: 70%"
29 echo "  - Min instances: 1"
30 echo "  - Max instances: 3"
31 echo "  - Cooldown period: 60 seconds"

```

Listing 3: Create a Managed Instance Group with auto-scaling

### Auto-Scaling Policy Parameters

Parameter	Value
Target CPU utilization	70%
Minimum instances	1
Maximum instances	3
Cooldown period	60 seconds

The screenshot shows the Google Cloud Console interface for the 'GCP-VM-assignment' project. The left sidebar lists various services, with 'Compute Engine' selected. The main panel displays 'VM instances' under the 'INSTANCES' tab. A table lists three instances: 'assignment-2-gcp-vm', 'autoscaling-vm01', and 'instance-group-1-b4b4'. Below the table, there are 'Related actions' such as 'Explore Backup and DR', 'View billing report', 'Monitor VMs', 'Explore VM logs', and 'Set up firewall rules'. A notification at the top suggests using zonal DNS.

Figure 4: Managed Instance Group with Initial VM Instance

## 4.3 Implementing Security Measures

### 4.3.1 Firewall Rules

```

1 #!/bin/bash
2 # Set up firewall rules to allow HTTP/HTTPS traffic only from a specific IP
3   (223.238.204.234/32)
4
5 PROJECT_ID="gcp-vm-assignment"
6 ALLOWED_IP="223.238.204.234/32"
7
8 gcloud config set project $PROJECT_ID
9
10 # Create firewall rule for HTTP traffic
11 gcloud compute firewall-rules create allow-http-custom \
12   --direction=INGRESS \
13   --priority=1000 \
14   --network=default \
15   --action=ALLOW \

```

```

15 --rules=tcp:80 \
16 --source-ranges=$ALLOWED_IP \
17 --target-tags=http-server
18
19 # Create firewall rule for HTTPS traffic
20 gcloud compute firewall-rules create allow-https-custom \
21 --direction=INGRESS \
22 --priority=1000 \
23 --network=default \
24 --action=ALLOW \
25 --rules=tcp:443 \
26 --source-ranges=$ALLOWED_IP \
27 --target-tags=https-server
28
29 echo "Firewall rules created successfully."
30 echo "HTTP (port 80) and HTTPS (port 443) traffic is now allowed only from
    IP: $ALLOWED_IP"

```

Listing 4: Set up firewall rules for restricted access

### Firewall Configuration

- Allows HTTP (port 80) and HTTPS (port 443) traffic only from the specified IP address (223.238.204.234/32)
- Targets instances with the "http-server" and "https-server" tags
- Sets a priority of 1000 for the rules

### 4.3.2 IAM Roles

```

1 #!/bin/bash
2 # Set up IAM roles to grant restricted access
3
4 PROJECT_ID="gcp-vm-assignment"
5 USER_EMAIL="user@example.com" # Replace with the email of the user
6
7 # Grant the Compute Viewer role to the specified user
8 gcloud projects add-iam-policy-binding $PROJECT_ID \
9 --member="user:$USER_EMAIL" \
10 --role="roles/compute.viewer"
11
12 echo "IAM role 'Compute Viewer' assigned to $USER_EMAIL"
13 echo "Note: This role provides read-only access to all compute resources"

```

Listing 5: Set up IAM roles for restricted access

### IAM Configuration

- Assigns the "Compute Viewer" role to a specified user
- Provides read-only access to all compute resources in the project

## 5 Testing

## 5.1 Testing Auto-Scaling

To test auto-scaling, we simulate high CPU load using the stress tool:

```
1 #!/bin/bash
2 # Stress test script to simulate CPU load on the VM instance
3 # This script is intended to be run on the VM (via SSH)
4
5 echo "Starting stress test to simulate high CPU load..."
6 echo "This test will help verify that auto-scaling triggers correctly."
7
8 # Install the stress tool if it is not already installed
9 if ! command -v stress &> /dev/null
10 then
11     echo "Installing stress tool..."
12     sudo apt-get update
13     sudo apt-get install -y stress
14 fi
15
16 # Run stress test on 2 CPU cores for 300 seconds (5 minutes)
17 echo "Running stress test on 2 CPU cores for 5 minutes..."
18 echo "Monitor the GCP Console to observe auto-scaling behavior."
19 stress --cpu 2 --timeout 300
20
21 echo "Stress test completed. Check the GCP Console to verify if auto-scaling
    was triggered."
```

Listing 6: Stress test script to simulate CPU load

## 5.2 Testing Firewall Rules

### 5.2.1 Testing with curl

```
1 #!/bin/bash
2 # Test HTTP response from your instance using curl
3
4 TARGET_IP="34.131.245.101"# Replace with your instance's public IP
5
6 echo "Testing HTTP connectivity to $TARGET_IP using curl..."
7 echo "This will verify if the web server is responding and firewall rules
   are properly configured."
8
9 # Run the curl command with verbose output
10 echo "Sending HTTP request to http://$TARGET_IP ..."
11 curl -v http://$TARGET_IP
12
13 echo ""
14 echo "If you received a successful HTTP response, you are accessing from an
   allowed IP."
15 echo "If the connection was refused or timed out, either your IP is not in
   the allowed list or the web server is not running."
```

Listing 7: Test HTTP response using curl

### 5.2.2 Testing with nmap

```
1 #!/bin/bash
2 # Test open port 80 on your instance using nmap
3
4 TARGET_IP="34.131.245.101"# Replace with your instance's public IP
5
6 echo "Testing HTTP port (80) accessibility on $TARGET_IP using nmap..."
7 echo "This will verify if the firewall rules are properly configured."
8
9 # Check if nmap is installed
10 if ! command -v nmap &> /dev/null
11 then
12     echo "nmap is not installed. Installing it now..."
```

```
13     sudo apt-get update
14     sudo apt-get install -y nmap
15 fi
16
17 # Run the port scan
18 echo "Scanning port 80 on $TARGET_IP ..."
19 nmap -p 80 $TARGET_IP
20
21 echo "If port 80 is reported as 'filtered' or 'closed' from an unauthorized
    IP, the firewall rule is working correctly."
```

Listing 8: Test port accessibility using nmap

### Firewall Testing Results

This script uses curl to test HTTP connectivity to the VM instance. When accessed from:

- **Allowed IP (223.238.204.234):** Successfully accessed the Apache web server content via HTTP (port 80)
- **Unauthorized IP (ping.eu):** Connection blocked, demonstrating effective firewall rule implementation

## 6 Deployment Manager Configuration

In addition to the scripts, we can use GCP Deployment Manager to define infrastructure as code. Below are the YAML configurations for firewall rules and IAM roles.

### 6.1 Firewall Rules Configuration

```
1 resources:
```

```
2 - name: allow-http-custom
3   type: compute.v1.firewall
4   properties:
5     description: "Allow HTTP traffic from allowed IP range "
6     network: global/networks/default
7     priority: 1000
8     sourceRanges:
9       - 223.238.204.234/32
10    allowed:
11      - IPProtocol: tcp
12        ports:
13          - "80"
14
15 - name: allow-https-custom
16   type: compute.v1.firewall
17   properties:
18     description: "Allow HTTPS traffic from allowed IP range "
19     network: global/networks/default
20     priority: 1000
21     sourceRanges:
22       - 223.238.204.234/32
23    allowed:
24      - IPProtocol: tcp
25        ports:
26          - "443"
```

Listing 9: YAML configuration for firewall rules

## 6.2 IAM Roles Configuration

```
1 resources:
2 - name: add-iam-policy-binding
3   type: gcp-types/cloudresourcemanager-v1:virtual.projects.iamMemberBinding
4   properties:
5     resource: projects/your-project-id
6     role: roles/compute.viewer
7     member: user: user@example.com
```

Listing 10: YAML configuration for IAM roles

## 7 Results and Verification

## 7.1 Auto-Scaling Results

### Auto-Scaling Behavior Observations

During testing, we observed the following behavior:

- When the CPU utilization exceeded 70%, new instances were added within approximately 2-3 minutes
- The number of instances scaled up to the maximum of 3 during peak load
- When the load decreased, instances were gradually removed after the cooldown period
- The MIG maintained at least 1 instance at all times

## 7.2 Security Verification

### Security Testing Results

The firewall rules were tested from both allowed and disallowed IP addresses:

- From the allowed IP (223.238.204.234), HTTP requests to port 80 received successful responses
- From disallowed IPs, HTTP requests were blocked, and port scans showed the port as "filtered"
- The IAM role assignment was verified by confirming that the user could view but not modify compute resources



## 8 Source Code Repository

### GitHub Repository

All source code, scripts, and configuration files used in this project are available in the following GitHub repository:

<https://github.com/AbhilashAgarwalIITJ/GCP-VM>

The repository includes:

- Scripts for VM creation and configuration
- Auto-scaling setup scripts
- Firewall and IAM configuration files
- Testing scripts and documentation
- Additional resources and references

## 9 Conclusion

This project successfully implemented a GCP Virtual Machine with auto-scaling capabilities and robust security measures. The solution provides:

- Dynamic scaling based on CPU utilization

- Restricted network access through firewall rules
- Controlled resource access through IAM roles
- Comprehensive testing procedures

The implementation demonstrates how GCP services can be combined to create a secure, scalable infrastructure that adapts to changing workload demands while maintaining strict access controls.

## References

- [1] Google Cloud Platform Documentation, <https://cloud.google.com/docs>
- [2] Google Compute Engine Documentation, <https://cloud.google.com/compute/docs>
- [3] Managed Instance Groups and Auto-scaling, <https://cloud.google.com/compute/docs/autoscaler>
- [4] VPC Firewall Rules, <https://cloud.google.com/vpc/docs/firewalls>
- [5] Identity and Access Management (IAM), <https://cloud.google.com/iam/docs>